

Patrycja Paradowska, 244952

Prophet and Neural Prophet

1. Prophet

Apply Prophet, SARIMA, and Holt-Winters' exponential smoothing to the Google trends vacation data set. The csv file is available on ePorta. Predict 12 months in the future. Don't forget to use monthly time series.

```
In [29]: import matplotlib.pyplot as plt
import pandas as pd
from fbprophet import Prophet
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.api import ExponentialSmoothing, Holt
from pmdarima import auto_arima
import warnings
from sklearn.metrics import mean_squared_error
from statsmodels.tools.eval_measures import rmse
warnings.filterwarnings("ignore")

In [30]: df=pd.read_csv('vacation.csv', parse_dates=True)
df
```

```
Out[30]:
```

Month	Num_Search_Vacation
0 2004-01-01	94
1 2004-02-01	89
2 2004-03-01	86
3 2004-04-01	79
4 2004-05-01	89
...	...
185 2019-06-01	58
186 2019-07-01	56
187 2019-08-01	45
188 2019-09-01	38
189 2019-10-01	37

190 rows x 2 columns

```
In [31]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 190 entries, 0 to 189
Data columns (total 2 columns):
 # Column          Non-Null Count  Dtype
---  --
 0 Month            190 non-null    object
 1 Num_Search_Vacation  190 non-null    int64
dtypes: int64(1), object(1)
memory usage: 3.1+ MB
```

```
In [32]: from statsmodels.tsa.stattools import adfuller
adfuller(df['Num_Search_Vacation'])

(-2.336474118969219,
 0.1604739704900124,
 177,
 ('1%': -3.467845319799907,
  '5%': -2.878011745497939,
  '10%': -2.575551186759871),
 885.8015084229178)
```

```
In [33]: df.columns = ['ds', 'y']

In [34]: df['ds'] = pd.to_datetime(df['ds'])

In [35]: df
```

```
Out[35]:
```

ds	y
0 2004-01-01	94
1 2004-02-01	89
2 2004-03-01	86
3 2004-04-01	79
4 2004-05-01	89
...	...
185 2019-06-01	58
186 2019-07-01	56
187 2019-08-01	45
188 2019-09-01	38
189 2019-10-01	37

190 rows x 2 columns

```
In [36]: df.set_index('ds').plot(figsize=(12, 8))

Out[36]: <AxesSubplot: xlabel='ds'>
```

```
In [37]: len(df)

Out[37]: 190
```

```
In [38]: train, test = df[0:(12)], df[(12):]
print('train size:', len(train), 'test size:', len(test))

Train set size: 178
Test set size: 12
```

```
In [41]: m = Prophet()
m.fit(train)
```

INFO:fbprophet:Disabling weekly seasonality. Run prophet with weekly_seasonality=True to override this.

INFO:fbprophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to override this.

fbprophet.forecaster.Prophet at 0x24a37814e88

```
In [42]: future = m.make_future_dataframe(periods = 12, freq = 'MS')

In [43]: future
```

```
Out[43]:
```

ds	y
0 2004-01-01	94
1 2004-02-01	89
2 2004-03-01	86
3 2004-04-01	79
4 2004-05-01	89
...	...
185 2019-06-01	58
186 2019-07-01	56
187 2019-08-01	45
188 2019-09-01	38
189 2019-10-01	37

190 rows x 1 columns

```
In [44]: forecast = m.predict(future)
forecast
```

```
Out[44]:
```

ds	trend	yhat_lower	yhat_upper	trend_lower	trend_upper	additive_terms	additive_terms_lower	additive_terms_upper	yhat
0 2004-01-01	82.502445	89.089835	98.766631	82.502445	82.502445	11.201315	11.201315	11.201315	11.201315
1 2004-02-01	82.205215	83.242934	92.932877	82.205215	82.205215	5.939019	5.939019	5.939019	5.939019
2 2004-03-01	82.205215	83.242934	92.932877	82.205215	82.205215	4.989236	4.989236	4.989236	4.989236
3 2004-04-01	81.829931	80.326051	86.587569	81.629931	81.629931	-0.388946	-0.388946	-0.388946	-0.388946
4 2004-05-01	81.822289	80.320191	86.582666	81.342289	81.342289	4.079325	4.079325	4.079325	4.079325
...
185 2019-06-01	51.829643	61.043549	70.623177	51.760563	51.895660	14.016379	14.016379	14.016379	14.016379
186 2019-07-01	51.828721	60.937527	70.307135	51.744242	51.914465	13.568549	13.568549	13.568549	13.568549
187 2019-08-01	51.827768	47.447702	57.133378	51.721502	51.934106	0.651789	0.651789	0.651789	0.651789
188 2019-09-01	51.826815	23.216854	41.796543	51.698284	51.957344	-14.999750	-14.999750	-14.999750	-14.999750
189 2019-10-01	51.825893	29.551407	39.396414	51.676797	51.974065	-17.541870	-17.541870	-17.541870	-17.541870

190 rows x 10 columns

```
In [45]: m.plot(forecast);

Out[45]:
```

```
In [46]: m.plot_components(forecast);

Out[46]:
```

SARIMA

```
In [9]: df2 = pd.read_csv('vacation.csv', index_col = 'Month', parse_dates = True)
```

```
In [10]: df2
```

```
Out[10]:
```

Month	Num_Search_Vacation
2004-01-01	94
2004-02-01	89
2004-03-01	86
2004-04-01	79
2004-05-01	89
...	...
2018-06-01	62
2018-07-01	61
2018-08-01	50
2018-09-01	38
2018-10-01	37

178 rows x 1 columns

```
In [14]: len(df2_test)
```

```
Out[14]: 12
```

```
In [17]: auto_arima(df2, seasonal=True, m=12, trace=True).summary()
```

Performing stepwise search to minimize aic

Model	AIC	BIC	Log Likelihood
ARIMA(1,1,1) (1,0,1) [12] intercept : AIC=984.654, Time=0.39 sec			
ARIMA(1,1,0) (0,0,0) [12] intercept : AIC=1402.455, Time=0.01 sec			
ARIMA(1,1,1) (0,1,1) [12] intercept : AIC=Inf, Time=0.20 sec			
ARIMA(1,1,1) (0,1,1) [12] intercept : AIC=1234.045, Time=0.17 sec			
ARIMA(1,1,0) (0,0,0) [12] intercept : AIC=1400.635, Time=0.01 sec			
ARIMA(2,1,2) (0,0,1) [12] intercept : AIC=Inf, Time=0.69 sec			
ARIMA(2,1,2) (0,1,2) [12] intercept : AIC=1006.205, Time=1.89 sec			
ARIMA(2,1,2) (1,0,2) [12] intercept : AIC=980.473, Time=1.72 sec			
ARIMA(2,1,2) (0,0,2) [12] intercept : AIC=Inf, Time=1.37 sec			
ARIMA(2,1,2) (0,0,2) [12] intercept : AIC=Inf, Time=1.86 sec			
ARIMA(1,1,2) (1,0,2) [12] intercept : AIC=979.429, Time=1.56 sec			
ARIMA(1,1,2) (0,0,2) [12] intercept : AIC=Inf, Time=1.23 sec			
ARIMA(1,1,2) (0,0,1) [12] intercept : AIC=Inf, Time=0.95 sec			
ARIMA(1,1,2) (0,0,2) [12] intercept : AIC=Inf, Time=1.86 sec			
ARIMA(1,1,2) (0,0,2) [12] intercept : AIC=1001.651, Time=0.57 sec			
ARIMA(1,1,1) (0,0,1) [12] intercept : AIC=979.483, Time=1.37 sec			
ARIMA(0,1,2) (1,0,2) [12] intercept : AIC=983.571, Time=1.00 sec			
ARIMA(1,1,1) (0,0,1) [12] intercept : AIC=976.697, Time=1.47 sec			
ARIMA(1,1,1) (0,0,1) [12] intercept : AIC=976.130, Time=0.19 sec			
ARIMA(1,1,1) (0,0,1) [12] intercept : AIC=978.359, Time=0.70 sec			
ARIMA(1,1,1) (0,0,1) [12] intercept : AIC=Inf, Time=1.64 sec			
ARIMA(1,1,1) (0,0,1) [12] intercept : AIC=976.130, Time=0.19 sec			
ARIMA(1,1,1) (0,0,1) [12] intercept : AIC=977.644, Time=1.61 sec			
ARIMA(0,1,1) (0,0,2) [12] intercept : AIC=990.860, Time=0.58 sec			
ARIMA(1,1,0) (1,0,2) [12] intercept : AIC=986.227, Time=0.49 sec			
ARIMA(2,1,1) (1,0,2) [12] intercept : AIC=Inf, Time=1.59 sec			
ARIMA(0,1,0) (1,0,2) [12] intercept : AIC=1001.651, Time=0.41 sec			
ARIMA(1,1,0) (1,0,2) [12] intercept : AIC=989.482, Time=0.61 sec			
ARIMA(1,1,1) (1,0,2) [12] intercept : AIC=975.306, Time=0.31 sec			
ARIMA(1,1,1) (0,0,1) [12] intercept : AIC=Inf, Time=1.39 sec			
ARIMA(1,1,1) (0,0,1) [12] intercept : AIC=1234.455, Time=0.11 sec			
ARIMA(1,1,1) (0,0,1) [12] intercept : AIC=976.130, Time=0.19 sec			
ARIMA(0,1,1) (1,0,2) [12] intercept : AIC=988.860, Time=0.51 sec			
ARIMA(1,1,2) (1,0,2) [12] intercept : AIC=994.227, Time=0.39 sec			
ARIMA(2,1,1) (1,0,2) [12] intercept : AIC=976.335, Time=0.91 sec			
ARIMA(1,0) (1,0,2) [12] intercept : AIC=999.651, Time=0.22 sec			
ARIMA(1,1,2) (1,0,2) [12] intercept : AIC=981.482, Time=0.56 sec			
ARIMA(2,1,0) (1,0,2) [12] intercept : AIC=997.692, Time=0.57 sec			
ARIMA(2,1,2) (1,0,2) [12] intercept : AIC=1000.524, Time=1.26 sec			

Best model: ARIMA(1,1,1) (1,0,2) [12]
Total fit time: 38.752 seconds

SARIMAX Results

Dep. Variable:	Num. Search_Vacation	No. Observations:	Log Likelihood:
Model:	SARIMAX(1, 1)x(1, 0, 1, 2) [12]	178	-481.653
Date:	Mon, 16 May 2022	AIC	975.306

Time:	13.0012	BIC	994.757
Sample:	0	HQIC	983.186
	-190		

Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.6628	0.058	11.367	0.000	0.549	-0.776
ma.L1	-0.9636	0.029	-32.979	0.000	-1.021	-0.907

	coef	std err	z	P> z	[0.025	0.975]
ar.S1.L2	0.9751	0.009	108.551	0.000	0.957	0.993
ma.S1.L2	-0.4194	0.071	-5.934	0.000	-0.558	-0.281

	coef	std err	z	P> z	[0.025	0.975]
ma.S1.L2	0.1426	0.069	2.074	0.038	0.008	-0.277
sigma2	8.1509	0.586	13.915	0.000	7.003	9.299

Ljung-Box (L1) (Q): 0.05 Jarque-Bera (JB): 59.16

	Prob(Q):	0.82	Prob(JB):	0.00
Heteroskedasticity (H):	1.59	Skew:	0.29	
Prob(H) (two-sided):	0.07	Kurtosis:	6.58	

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [18]: model=Sarimax = SARIMAX(df2_train, order=(1,1,1), seasonal_order=(1,0,2,12), enforce_invertibility=False, enforce_stationarity=True)
fitted = model.fit(df2_train)
fitted.summary()
```

```
Out[18]:
```

Dep. Variable:	Num. Search_Vacation	No. Observations:	Log Likelihood:
Model:	SARIMAX(1, 1)x(1, 0, 1, 2) [12]	178	-368.628
Date:	Mon, 16 May 2022	AIC	749.256

Time:	13.0029	BIC <th>767.359</th>	767.359
Sample:	01-01-2004	HQIC	756.610
	-10-01-2018		

Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.5774	0.108	5.329	0.000	0.365	0.790
ma.L1	-0.9143	0.073	-12.580	0.000	-1.057	-0.772

	coef	std err	z	P> z	[0.025	0.975]
ar.S1.L2	0.9339	0.017	54.525	0.000	0.900	0.967
ma.S1.L2	-0.4841	0.087	-5.565	0.000	-0.655	-0.314

	coef	std err	z	P> z	[0.025	0.975]
ma.S1.L2	0.0950	0.078	1.224	0.221	-0.057	0.249
sigma2	7.5160	0.547	13.743	0.000	6.445	8.587

Ljung-Box (L1) (Q): 0.08 Jarque-Bera (JB): 88.94


	Prob(Q):	0.78	Prob(JB):	0.00
Heteroskedasticity (H):	2.47	Skew:	0.60	
Prob(H) (two-sided):	0.00	Kurtosis:	6.56	

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [20]: pred_Sarimax=fitted.predict(start=start, end=end, dynamic=False)

In [21]: fitted.plot_diagnostic(figsize=(12,6));
```




```
In [22]: ax = df2_train.plot(figsize=(12,6))
df2_train.plot(ax=ax)
pred_Sarimax.plot(ax=ax)
plt.legend(['Train', 'Test', 'SARIMA forecast'])

Out[22]: <matplotlib.legend.Legend at 0x24a3a162808>
```



```
In [24]: ax = df2_test.plot(figsize=(12,6))
pred_Sarimax.plot(ax=ax)
plt.legend(['Test', 'SARIMA forecast'])

Out[24]: <matplotlib.legend.Legend at 0x24a3a20c4d0>
```



Holt-Winters' exponential smoothing

```
In [25]: add_add = ExponentialSmoothing(df2_train, trend='add', seasonal='add', seasonal_periods=12).fit()
add_mult = ExponentialSmoothing(df2_train, trend='add', seasonal='mul', seasonal_periods=12).fit()
add_mult = ExponentialSmoothing(df2_train, trend='mul', seasonal='add', seasonal_periods=12).fit()
add_mult = ExponentialSmoothing(df2_train, trend='mul', seasonal='mul', seasonal_periods=12).fit()
```

```
In [26]: add_add_forecast = add_add.forecast(len(df2_test))
add_mult_forecast = add_mult.forecast(len(df2_test))
add_mult_forecast = add_mult.forecast(len(df2_test))
add_mult_forecast = add_mult.forecast(len(df2_test))
```

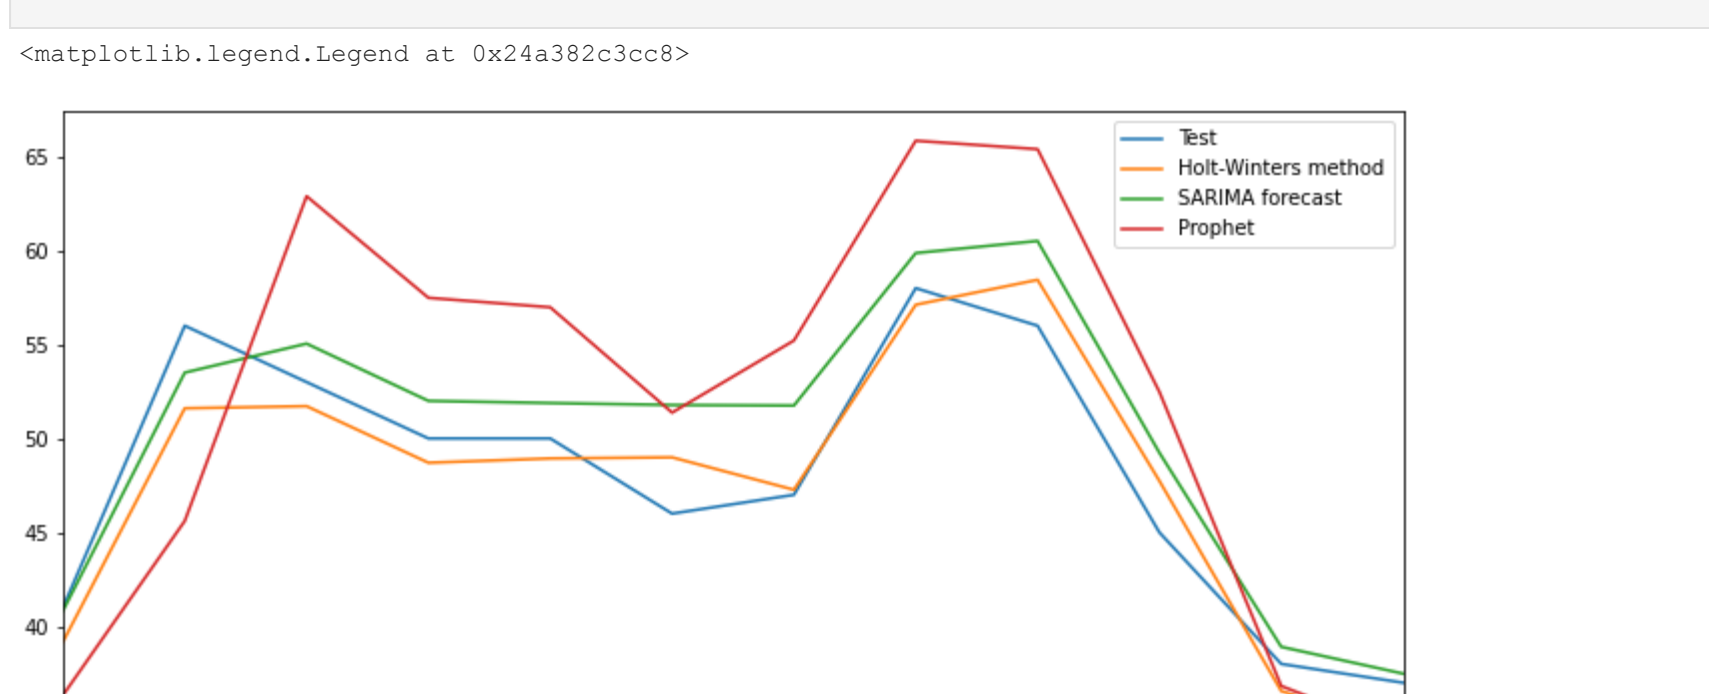
```
In [27]: mean_squared_errors = [
    (mean_squared_error(df2_test, add_add_forecast)),
    (mean_squared_error(df2_test, pred_Sarimax)),
    (mean_squared_error(df2_test, add_mult_forecast)),
    (mean_squared_error(df2_test, add_mult_forecast))
]

mean_squared_errors
```

```
Out[27]: [4.952623407325601, 12.614135770163188, 6.2756507014344765, 13.50500046631454]
```

```
In [53]: ax = df2_test.plot(figsize=(16, 10), label = "TEST")
add_add_forecast.plot(ax=ax, label="add_add_forecast")
add_mult_forecast.plot(ax=ax, label="add_mult_forecast")
add_mult_forecast.plot(ax=ax, label="add_mult_forecast")
add_mult_forecast.plot(ax=ax, label="add_mult_forecast")
ax.legend(['TEST', 'add_add_forecast', 'add_mult_forecast', 'add_mult_forecast'])

Out[53]: <matplotlib.legend.Legend at 0x24a3820c388>
```

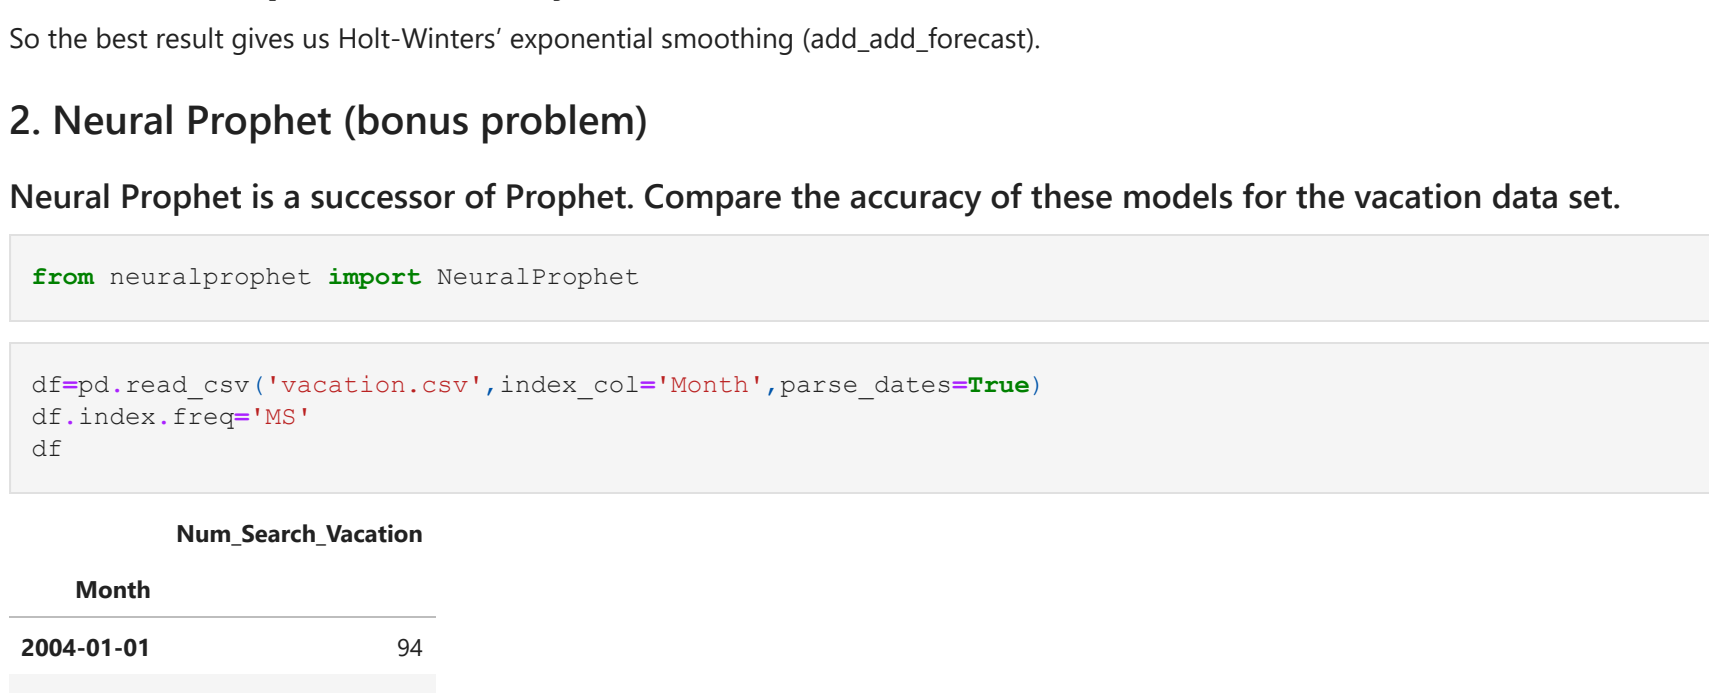


So the best result gives us add_add_forecast (4.952623407325601).

Compare the accuracy of the results:

```
In [56]: prop = pd.Series(data=results["yhcat"])
index = prop.to_frame()
prop_df = prop.to_frame().iloc[0:12]
ax = df2_test.plot(figsize=(12,6))
add_add_forecast.plot(ax=ax)
pred_Sarimax.plot(ax=ax)
pred_df["yhcat"].plot(ax=ax)
plt.legend(['Test', 'Holt-Winters method', 'SARIMA forecast', 'Prophet'])

Out[56]: <matplotlib.legend.Legend at 0x24a3820c3cc>
```



```
In [57]: error=mean_squared_error(df2_test,prop_df['yhcat'][-12:])
print("PROPHET: ", error)
error=mean_squared_error(df2_test, pred_Sarimax)
print("SARIMA: ", error)
error=mean_squared_error(df2_test, add_add_forecast)
print("Holt-Winters' exponential smoothing: ", error)
```

PROPHET: 53.60896752764628
SARIMA: 9.739922997734

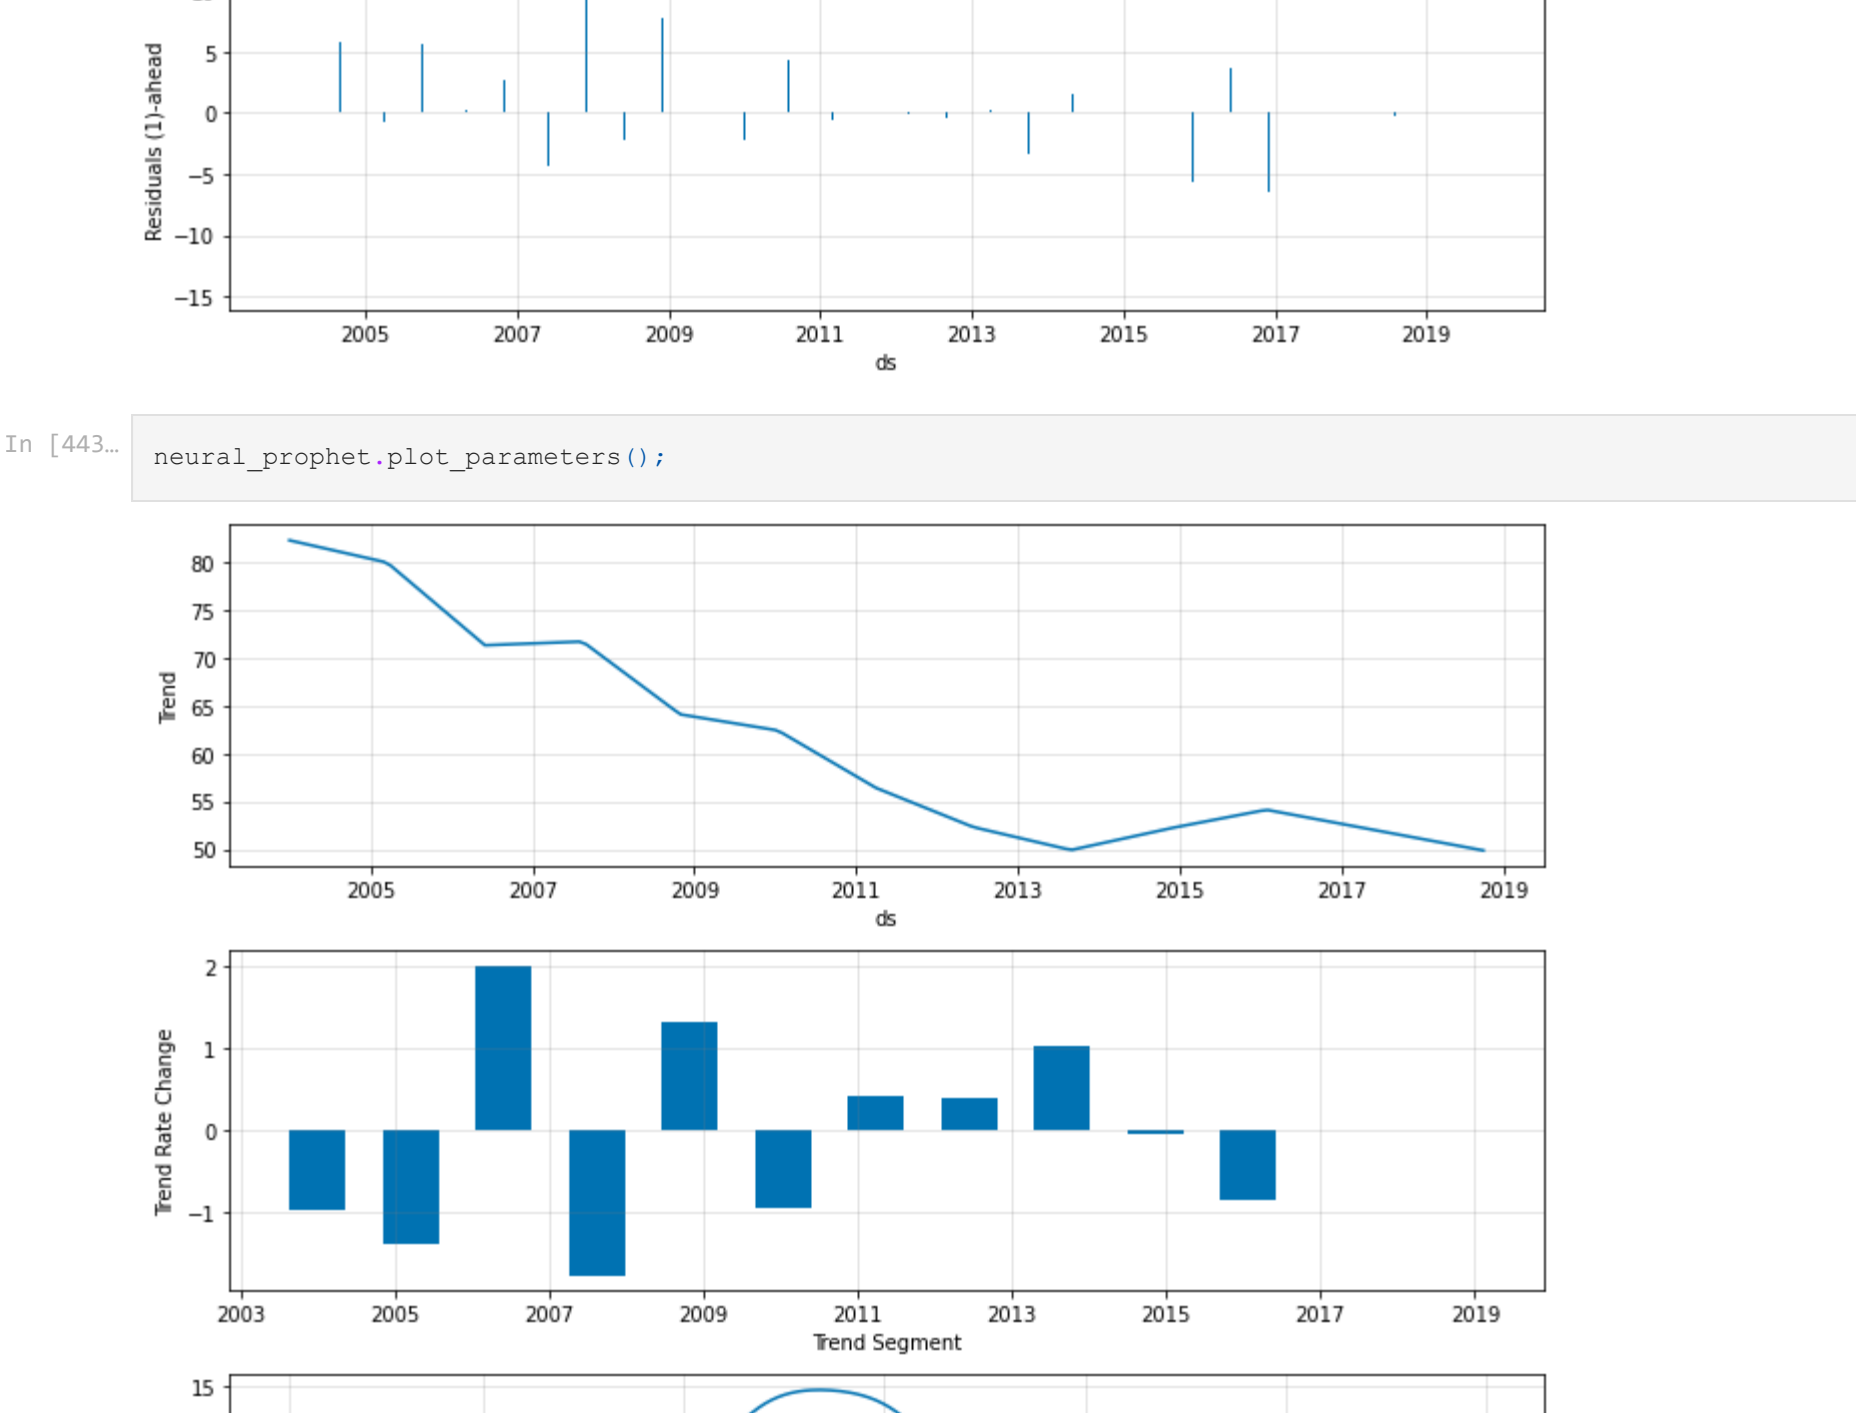

```
In [440]: neural_prophet = NeuralProphet()
neural_prophet.fit(train_data, freq='MS')
future = neural_prophet.make_future_dataframe(train_data, periods=12, n_historic_predictions=len(train_data))
neural_prophet_forecast = neural_prophet.predict(future)

INFO - (NP.df_utils.infer_frequency) - Major frequency MS corresponds to 91.011% of the data.
INFO:NP.df_utils:Major frequency MS corresponds to 91.011% of the data.
INFO - (NP.df_utils.infer_frequency) - Defined frequency is equal to major frequency - MS
INFO:NP.df_utils:Defined frequency is equal to major frequency - MS
INFO - (NP.config.init_data_params) - Setting normalization to global as only one dataframe provided for training.
INFO:NP.config:Setting normalization to global as only one dataframe provided for training.
INFO - (NP.utils.set_auto_seasonalities) - Disabling weekly seasonality. Run NeuralProphet with weekly_seasonality=True to override this.
INFO:NP.utils:Disabling weekly seasonality. Run NeuralProphet with weekly_seasonality=True to override this.
INFO - (NP.utils.set_auto_seasonalities) - Disabling daily seasonality. Run NeuralProphet with daily_seasonality=True to override this.
INFO:NP.utils:Disabling daily seasonality. Run NeuralProphet with daily_seasonality=True to override this.
INFO - (NP.config.set_auto_batch_epoch) - Auto-set batch_size to 16
INFO:NP.config:Auto-set batch_size to 16
INFO - (NP.config.set_auto_batch_epoch) - Auto-set epochs to 388
INFO:NP.config:Auto-set epochs to 388

INFO - (NP.utils.torch.lr_range_test) - lr-range-test results: steep: 1.13E+01, min: 9.01E-01
INFO:NP.utils.torch:lr-range-test results: steep: 1.13E+01, min: 9.01E-01

INFO - (NP.utils.torch.lr_range_test) - lr-range-test results: steep: 9.39E+02, min: 1.58E+00
INFO:NP.utils.torch:lr-range-test results: steep: 9.39E+02, min: 1.58E+00
INFO - (NP.forecaster._init_train_loader) - lr-range-test selected learning rate: 1.02E-01
INFO:NP.forecaster:lr-range-test selected learning rate: 1.02E-01
Epoch[388/388]: 100% 388/388 (00:07<00:00, 50.36it/s, SmoothLoss=0.00231, MAB=2.87, RMSE=3.66, Rsgloss=0)

INFO - (NP.df_utils.infer_frequency) - Major frequency MS corresponds to 91.011% of the data.
INFO:NP.df_utils:Major frequency MS corresponds to 91.011% of the data.
INFO - (NP.df_utils.infer_frequency) - Defined frequency is equal to major frequency - MS
INFO:NP.df_utils:Defined frequency is equal to major frequency - MS
INFO - (NP.df_utils.infer_frequency) - Major frequency MS corresponds to 91.053% of the data.
INFO:NP.df_utils:Major frequency MS corresponds to 91.053% of the data.
INFO - (NP.df_utils.infer_frequency) - Defined frequency is equal to major frequency - MS
INFO:NP.df_utils:Defined frequency is equal to major frequency - MS
INFO - (NP.df_utils.infer_frequency) - Major frequency MS corresponds to 91.053% of the data.
INFO:NP.df_utils:Major frequency MS corresponds to 91.053% of the data.
INFO - (NP.df_utils.infer_frequency) - Defined frequency is equal to major frequency - MS
INFO:NP.df_utils:Defined frequency is equal to major frequency - MS
INFO - (NP.df_utils.infer_frequency) - Major frequency MS corresponds to 91.053% of the data.
INFO:NP.df_utils:Major frequency MS corresponds to 91.053% of the data.
INFO - (NP.df_utils.infer_frequency) - Defined frequency is equal to major frequency - MS
INFO:NP.df_utils:Defined frequency is equal to major frequency - MS
```



```
In [444]: np_df = pd.Series(data=neural_prophet_forecast.loc[:, 'yhat1']).to_frame()
np_df = np_df.set_index(pd.date_range(start = "2004-01-01", end = "2019-11-01", freq = "M")).iloc[-12:]

In [445]: ax = test_data['y'].plot(figsize=(16, 10))
np_df.plot(ax=ax)
prop_df.plot(ax=ax)
plt.legend(['Test', 'Neural prophet', 'Prophet'])

Out[445]: <matplotlib.legend.Legend at 0x18f46b62b88>
```



```
In [446]: print('MSE comparison:')
print(f'Prophet: {mean_squared_error(test_data['y'], prop_df['yhat1']):.4f}')
print(f'NeuralProphet: {mean_squared_error(test_data['y'], np_df['yhat1']):.4f}')

MSE comparison:
Prophet: 53.6090
NeuralProphet: 31.3549

So we see that Neural Prophet gives better results than Prophet.
```